

How the four papers build toward one goal:
memory-optimal kernels for the complete transformer block, validated on real hardware

