

Validating Memory-Optimal Transformer Kernels on Real Hardware: From Formal Derivation to Measured Performance Across Two HPC Clusters

Anonymous Author(s)

Submission under double-anonymous review

Abstract—A companion series of papers formally derives memory-optimal cost functions for the components of a transformer block — attention’s forward pass, backward pass, fused forward+backward, inference-time decode, and the complete block — each verified to machine precision against PyTorch autograd, but never checked against real hardware. This paper reports that validation. Cost-function predictions are checked against measured performance on two HPC clusters, across CPU and GPU, using the Mathematics of Arrays (MoA) formalism’s shape/index vocabulary (ρ , ψ , ι) both to derive each kernel’s machine-specific realization (an Operational Normal Form, ONF) from its hardware-independent specification (a Denotational Normal Form, DNF) and to diagnose gaps when prediction and measurement diverge. Three results stand out. First, we identify and fix a real GPU performance regression: fusing forward and backward passes — proven to avoid materializing an $O(n^2)$ intermediate array — initially ran *slower* than the naive alternative on GPU, due to atomic-memory contention in the generated kernel. Profiling confirmed the mechanism to four decimal places ($2.0000\times$ more atomic instructions than a structurally related kernel), and a targeted restructuring of the ONF, guided directly by that diagnosis, reversed the result completely, yielding up to $2.5\times$ real speedup. Second, we show the identical derivation produces markedly different real costs on different machine topologies — a $535\times$ NUMA-locality penalty on one cluster’s architecture versus under $3\times$ oversubscription cost on another’s — demonstrating that optimal deployment is a function of the target machine’s own array structure, not a fixed property of the algorithm. Third, we report a genuine, only partially resolved anomaly: identical denotational computations run faster in C than Fortran on CPU but faster in Fortran than C on GPU; we narrow this reversal to one dominant kernel and one stall mechanism without fully explaining its compiler-level cause. Together, this evidence supports treating hardware-specific optimization as a routine, targeted rewrite of a kernel’s machine-specific realization alone, with its underlying formal derivation remaining fixed, verified, and reusable across hardware generations.

I. INTRODUCTION

Formally derived cost functions are only useful in practice if they predict what real hardware actually does. A companion series of papers derives, from first principles in the Mathematics of Arrays (MoA) formalism, memory-optimal cost functions for every component of a transformer block, each verified to machine precision against PyTorch autograd. That verification, however, was correctness-only: conducted

at a single small problem size, on one CPU core or via a GPU compiler’s host-fallback path. It establishes that the derived kernels compute the right answer; it says nothing about whether their *cost predictions* hold at scale, on real multi-core CPUs and real GPUs. This paper reports that missing validation.

MoA’s five primitives — ψ (indexing), ι (index generation), ρ (shape), γ (the translator from a denotational to an operational form), and rav (ravel) — let a kernel’s output be specified purely in terms of shape and index relationships, with no commitment to memory layout or access order: a *Denotational Normal Form* (DNF). γ translates a DNF into an *Operational Normal Form* (ONF) that commits to loop order, tiling, and which intermediates are recomputed versus stored. Crucially, γ ’s output remains expressed in the same shape/index vocabulary as the DNF — because on real hardware, an array’s shape and index structure *is* its memory layout and access pattern. This is the sense in which MoA treats the target machine itself as an array: DRAM address space, cache lines, and thread-to-data assignment are describable in $\rho/\psi/\iota$ terms, so a γ -derived ONF is a literal, checkable claim about what bytes move where on real silicon, not a metaphorical implementation guide. This paper’s method is a direct test of that claim: if measured DRAM traffic, atomic-instruction counts, or thread scaling match what an ONF predicts, the machine-as-array claim holds; where they diverge, the gap should be diagnosable in the same vocabulary the prediction was stated in, rather than requiring ad hoc explanation.

We validate this across four kernel families (attention forward pass, backward pass, fused forward+backward, inference-time decode, and the complete transformer block), on two independent HPC clusters, across both CPU and GPU targets. Section II describes the experimental setup. Section III reports a real GPU regression, its diagnosis, and its resolution — the clearest case in this paper of a formally-guided fix confirmed by measurement. Section IV reports a cross-machine finding showing that optimal deployment is a function of target-machine topology, not a fixed algorithmic property. Section V reports a genuine, only partially resolved compiler-level anomaly. Additional validation results are summarized in Section VI. Section VII discusses implications for treating

hardware-specific optimization as a routine, targeted activity distinct from formal correctness.

II. BACKGROUND AND METHODOLOGY

A. Two clusters, real allocations

Experiments ran on two NSF ACCESS-allocated HPC clusters: Cluster A (128-core CPU nodes, NVIDIA A100/H100 GPU nodes) and Cluster B (128-core CPU nodes with 8 NUMA domains, NVIDIA A100/A40/H200 GPU nodes). All GPU kernels are compiled with an OpenACC-capable compiler targeting NVIDIA hardware; CPU kernels use OpenMP. Each kernel is independently verified for correctness (machine-precision agreement against a PyTorch or serial-CPU reference) prior to any performance run reported here; this paper reports timing, not correctness, except where a correctness issue is itself the finding.

B. A methodological artifact worth reporting: GPU cold-start dominance

The first real GPU measurement obtained in this effort showed wall-clock time essentially flat across a $1024\times$ range in sequence length — inconsistent with the kernel’s $O(n)$ cost. The cause: the original benchmark harness launched one process per measured point, and CUDA context creation plus JIT compilation cost approximately 300ms per launch, completely dominating true kernel time (sub-millisecond even at $n = 10^6$). Every GPU benchmark in this project shared this flaw. The fix, applied uniformly: sweep the entire n range within one process, paying context-init cost once via an explicit untimed warmup. We report this because it generalizes: a naive per-call GPU timing harness will systematically overstate cost by a roughly constant offset at every problem size, obscuring the true asymptotic complexity of any sub-millisecond kernel unless that constant is isolated first.

III. CASE STUDY: GPU ATOMIC CONTENTION

Fusing attention’s forward and backward passes into a single kernel is proven, by construction, to avoid materializing an $n \times n$ intermediate score matrix that the naive (unfused) alternative must write to and read back from DRAM — an $O(n^2)$ -scaling cost. On CPU, this predicted advantage is confirmed directly: at $N = 4096$, fusion is 28% faster than naive at the maximum tested thread count, with the advantage growing monotonically as thread count increases and memory bandwidth contention makes the naive path’s extra traffic proportionally more expensive.

A. An unexpected regression

On GPU, the identical fusion is *slower* than naive at every tested problem size, on two GPU architectures: up to $1.84\times$ slower on one, holding near $1.5\times$ slower on the other regardless of n . This is the opposite of the CPU result, for a fusion whose memory-traffic advantage is proven independent of target architecture.

TABLE I
FUSED-KERNEL TIME RELATIVE TO NAIVE (FORWARD+BACKWARD SEPARATELY), BEFORE AND AFTER THE ATOMIC-ELIMINATION FIX, TWO REAL GPU ARCHITECTURES.

	Before fix	After fix
GPU architecture A	1.14–1.84× slower	0.52–0.76× (faster)
GPU architecture B	1.34–1.53× slower	0.40–0.92× (faster)

B. Diagnosis in the ONF’s own vocabulary

Profiling both kernels with a hardware performance-counter tool revealed the mechanism. The fused kernel’s γ -derived ONF for GPU requires two atomic-accumulation sites (for two distinct gradient arrays); a structurally related unfused kernel’s ONF requires only one, in the same form. This is a specific, countable, testable prediction, made from the ONF’s structure alone, before any measurement. Extracting atomic-instruction counts confirmed it precisely: the fused kernel executes exactly $2.0000\times$ the atomic instructions of the related kernel, and device time tracks almost identically ($1.9716\times$). Per-kernel occupancy analysis adds a further finding: the atomic-using computation is both the slowest component of the unfused kernel (85% of its total time) and its least-occupied (14.2% versus 44–48% for atomic-free components) — already a bottleneck before fusion, which then inherits and doubles it.

C. Resolution: restructuring the ONF, not the DNF

The fused kernel was restructured into multiple passes, splitting the monolithic kernel so that one gradient array’s accumulation could be parallelized without an atomic, eliminating that site while leaving the correctness-verified DNF completely unchanged. The fix was verified to exact bit-level agreement with the original implementation before any hardware run.

The result is not a partial improvement but a complete reversal on both architectures (Table I): fused is now faster than naive at every tested size, up to $2\times$ faster on one architecture and $2.5\times$ on the other, with the improvement growing as problem size increases — consistent with atomic contention mattering more as more threads compete for the same memory locations. This closes the central empirical question for this kernel pair: fusion is the right choice on both CPU and GPU once the GPU implementation’s own atomic-contention cost, orthogonal to the DRAM-traffic argument the formal derivation makes, is addressed.

D. A follow-up finding: not every guided fix generalizes cleanly

The remaining atomic site (in a second gradient computation) was deliberately left in place for the fix above. Removing it too, via the same technique, was tested as a natural follow-up. The result is genuinely mixed: for the unfused kernel alone, removing this second atomic helps uniformly (19–44% faster at every tested size); for the fused kernel, it helps at small and medium problem sizes but is measurably *slower* at large sizes (1.04 – $1.32\times$ slower). Profiling the mechanism directly

TABLE II
DECODE KERNEL: TIME RATIO, FULL THREAD COUNT VERSUS HALF, BY
SEQUENCE LENGTH n , CLUSTER B.

n	Full vs. half threads
1024	$535\times$ slower
8192	$183\times$ slower
65536	$0.6\times$ (faster)

(rather than speculating) found the cause: the new, atomic-free pass re-reads an intermediate array from DRAM a second time, at almost exactly the DRAM traffic of the pass that first computes it, adding real memory traffic that outweighs the removed atomic’s cost at large n . Despite this, the fused kernel still beats naive at every tested size overall, since the naive baseline’s own unfused component improves from the same fix. The lesson: a fix confirmed for one kernel does not transfer to a structurally different kernel without separately checking it, even when the underlying technique and the diagnosis method are identical.

IV. CASE STUDY: CROSS-MACHINE TOPOLOGY EFFECTS

If an ONF’s cost prediction genuinely depends on a target machine’s own array structure, as Section III’s framing claims, then the identical derivation should produce different real costs on machines with different memory topologies — and that difference should be explainable in topology terms, not shrugged off as noise.

A. A large, real penalty on one cluster

The inference-time decode kernel, run at its maximum thread count versus half that count on Cluster B’s CPU nodes, shows a striking, size-dependent penalty (Table II): at small sequence lengths, using twice as many threads is over $500\times$ slower, not merely inefficient.

Cluster B’s CPU nodes are confirmed, via direct topology inspection, to subdivide each socket into four NUMA domains (eight total across two sockets). The full thread count spans all eight domains and both sockets’ cross-socket interconnect; at small n , decode’s per-token work is minimal, so threads scattered across all eight domains pay full remote-memory latency for cache accesses while doing almost no compute to amortize that cost against. The penalty resolves as n grows large enough for genuine computation to dominate — by the largest tested size, the ratio is under $1.3\times$, not two orders of magnitude.

B. A structurally different penalty on another cluster

The identical, unmodified binary was also run on Cluster A, whose CPU nodes are a single-socket, 32-core design with simultaneous multithreading disabled. Here, the analogous oversubscription comparison (double the physical core count) shows a maximum penalty of under $3\times$ — roughly two orders of magnitude smaller than Cluster B’s worst case — and, at large n , *reverses into a genuine benefit* (up to $1.8\times$ faster), a textbook signature of latency-hiding oversubscription rather

than NUMA-domain remote-access cost. The two clusters’ penalties differ by two orders of magnitude because they are different mechanisms, confirmed by each cluster’s own topology, not the same mechanism at different severity.

C. Implication for deployment

The correct deployment choice (here, thread count) is a function of the target machine’s own array structure, not a fixed constant an algorithm’s specification could specify once. A single hardcoded default, chosen for one machine and never revisited, is not marginally suboptimal on a different machine — it can be actively harmful by a factor reaching into the hundreds.

V. CASE STUDY: A COMPILER-LEVEL ANOMALY

The complete transformer block was implemented in both C and Fortran, compiled through the same OpenACC toolchain, both verified to machine precision against the identical reference. On CPU, C is consistently faster than Fortran for the gated-MLP kernel, by $1.2\text{--}1.85\times$ across eleven tested sizes. On GPU, the same denotational computation reverses: Fortran is faster than C, by $1.45\text{--}2.81\times$. Both directions are real, repeated findings, not noise on either side.

A. A first attempt: instruction-level comparison

Since both languages compile to the identical ONF by construction, whatever explains the reversal lives in code generation. Extracting and categorizing compiled GPU machine code by instruction type tested the simplest hypothesis — that the faster language generates *fewer* instructions. It does not: Fortran generates *more* instructions than C in every category (11.4% more in total) despite running faster. This rules out a simple instruction-count explanation without providing a replacement one.

B. A follow-up: where do cycles actually go

Warp stall-reason profiling went further, asking not what instructions exist but where cycles are spent. One sub-kernel — a gated-MLP backward pass — dominates total runtime for both languages (88% of C’s profiled time, 65% of Fortran’s), making it the one that actually determines the aggregate result. For that kernel specifically, C runs $3.17\times$ slower than Fortran, and C’s warps stall on global-memory latency $3.35\times$ more often — closely matching the timing gap. This is a real, positive, quantified finding, not merely a ruled-out explanation, though it does not identify *why* C’s generated code for this one kernel produces worse memory-latency behavior than Fortran’s; that would require inspecting the generated memory-access pattern directly, which we have not yet done. The question has narrowed from an unexplained whole-pipeline reversal to one kernel and one stall mechanism.

VI. ADDITIONAL VALIDATION RESULTS

Beyond the three case studies above, direct profiler-based validation against closed-form DRAM-traffic predictions was performed for two further kernels. For the forward-pass kernel, measured DRAM traffic exceeds the naive element-count

prediction by $2.01\times$, traced via memory-coalescing efficiency metrics to a specific, addressable stride pattern (global stores achieving only 29% of theoretical coalescing efficiency); the underlying algorithm is memory-optimal by construction, and the gap is a code-generation issue, not a derivation issue. For the complete transformer block, the analogous gap is $1.50\times$; applying the same coalescing diagnosis found real, measured inefficiency directionally consistent with the gap but insufficient to arithmetically reconstruct it in full, indicating an additional, unidentified contributing factor beyond coalescing alone. The complete block’s real GPU scaling independently confirms its predicted $O(n^2)$ growth cleanly across three doublings of problem size. Finally, a growth-rate anomaly at one specific problem size in the complete block’s CPU sweep, initially ambiguous between a genuine cache-capacity effect and single-shot measurement noise, was confirmed genuine (not noise) by re-running with proper repeated-averaging: the anomaly persisted and slightly grew under averaging rather than smoothing toward the theoretical value, which noise would be expected to do.

VII. DISCUSSION

Across every case study in this paper, when a real-hardware measurement diverged from a formal prediction, the productive next step was to look for the cause inside the ONF’s own shape/index structure before looking anywhere else — and in every case but one (Section V), that search found a specific, fixable, correctly-diagnosed cause, subsequently confirmed by a second round of measurement rather than assumed. We take this as evidence for a broader claim about how AI systems might be built and maintained as hardware continues to change: separating a hardware-independent specification from its machine-specific realization is not merely a theoretical convenience but a practical scaling strategy. Conventional practice re-tunes kernels for each new architecture by profiling, guessing, and iterating without a formal predictive model to check the guess against — expensive in engineer-hours, and every un-retuned kernel in the meantime silently burns extra compute and energy moving more memory traffic than necessary. The fixes in this paper did not work that way: each was found by reading the ONF’s own structure, and confirmed against real measurement rather than merely hoped to generalize. The formally-derived cost functions, verified once against a correctness reference and now checked against real-machine outcomes across two independent clusters, are the reusable part; each machine’s realization is the part that legitimately changes, and changes with a specific, checkable target each time rather than an open-ended search.

Not every question resolved cleanly. Section V’s compiler-level reversal remains only partially explained, and we report that honestly rather than either overclaiming a full explanation or omitting the finding. Whether the pattern demonstrated here generalizes beyond the kernels and clusters tested in this paper, to the wider space of AI workloads and the full diversity of production hardware, remains a research direction rather than a conclusion this paper reaches on its own — but the pattern

held, without exception, on every kernel family it was tried on here.

VIII. CONCLUSION

We asked a narrow empirical question — do memory-optimal cost functions derived formally from array algebra actually predict real hardware behavior? — and answered it across four kernel families, two independent clusters, and both CPU and GPU targets. Where predictions held, they held to within a diagnosable, mechanistically understood margin; where they did not initially hold, the gap itself pointed to its own fix in the same formal vocabulary the prediction was stated in, and that fix was confirmed, not assumed. This supports treating hardware-specific optimization as a routine, targeted rewrite of a kernel’s machine-specific realization alone, with the underlying formal derivation remaining fixed, verified, and reusable as hardware continues to change.

ACKNOWLEDGMENTS

This work used NSF ACCESS-allocated HPC resources. Portions of this manuscript’s text were drafted with the assistance of an AI language model (Anthropic Claude); all technical content, experimental results, and conclusions were verified by the authors against source data prior to submission.

REFERENCES

- [1] Anonymized, *Attention at the Theoretical Minimum I: Forward Pass*, arXiv preprint, 2026.
- [2] Anonymized, *Attention at the Theoretical Minimum II: Backward Pass and Fused Forward+Backward*, HAL preprint, 2026.
- [3] Anonymized, *MoA-Structured Decode Attention: DNF Derivation, KV-Cache Accumulation, GQA/MQA, and OpenACC Kernel*, arXiv preprint, 2026.
- [4] Anonymized, *Attention at the Theoretical Minimum: A Unified Treatment of the Transformer Block*, submitted to IEEE Transactions on Neural Networks and Learning Systems, 2026.
- [5] Anonymized, *A Mathematics of Arrays*, Ph.D. dissertation, 1988.